

Length generalization, measured in the weights

A controlled study of six positional encodings on a small arithmetic transformer

Sanjay Ananthanarayan Independent research · github.com/sananthanarayan/jax-transformer

Version 1.0 · May 2026 · Technical report

Keywords: length generalization, positional encoding, mechanistic interpretability, arithmetic transformers, NoPE, RoPE, Abacus embeddings

Abstract

I trained a 10.6M-parameter decoder-only transformer on 3-digit addition and tested its ability to generalize to operand sizes never seen at training time (1–6 digits), comparing six positional-encoding variants over three seeds on the same architecture and schedule: learned absolute, learned absolute with reversed answers, no positional encoding (NoPE), rotary (RoPE), Abacus-style place-value embeddings, and Abacus trained with a length curriculum. Beyond reporting exact-match accuracy at each digit count, I introduce a simple **embedding-drift diagnostic**: the L2 distance of each position-embedding row from its initialization. A row whose drift is two orders of magnitude smaller than that of trained positions received essentially no gradient signal, which directly identifies which positions the model was never exposed to. The finding is that at this scale **no mechanism extrapolates by itself** — NoPE, RoPE, and clean Abacus all cliff to 0.0 ± 0.0 accuracy at the training-distribution boundary, bit-exact across seeds. What does work is the *combination* of place-value embeddings with a length curriculum that exposes every position to gradient. The drift diagnostic confirms this directly in the weights: the cliff position is identical across all three seeds, and the curriculum variant has non-zero drift exactly one position past where clean Abacus stops.

Reproducibility. All numbers in §6 are mean $\pm$ standard deviation over 3 random seeds (full sweep logged in `results/sweep.json`, ~3.4 hours of CPU). Figures show the mean with $\pm 1\sigma$ shading. The entire pipeline reproduces from one command: `make figures`.

1. Introduction

Length generalization on arithmetic is a sharp probe for what a transformer learns about position. The task is small enough that controlled experiments fit comfortably on a single GPU in minutes; the rules of the task are unambiguous; and yet the model has to *systematically* extend a pattern to inputs longer than anything it saw. This is the setting in which length-generalization failures of standard transformers were first cleanly demonstrated [Anil et al. 2022](#), and the setting in which several recent proposals (NoPE [Kazemnejad et](#)

al. 2023, Abacus embeddings McLeish et al. 2024, scratchpad reasoning Nye et al. 2021) have been benchmarked.

The setup here is deliberately minimal:

- A 10.6M-parameter decoder-only transformer with pre-norm, GeLU MLPs, RMSNorm, and causal self-attention (no fancy attention variants).
- A 15-token character-level vocabulary (digits, space, +, *, =, <pad>).
- Synthetic data: "{a} + {b} = {result}", optionally with the answer reversed ("579" written as "975" so carries flow with the causal direction Nogueira et al. 2021).
- Training on the full 10^6 pairs of 3-digit operands, with held-out OOD evaluation at 1, 2, 3, 4, 5, and 6 digits.

The contributions of this report are:

1. **A clean six-variant comparison** of positional encodings on the same architecture, training schedule, and data. Most prior work compares fewer variants or varies more than one factor at a time.
2. **The embedding-drift diagnostic**, which turns the abstract claim “positions never seen at training stay at random init” into a measurable per-row L2 number that anyone can read off the model’s weights.
3. **A length-curriculum control** for the Abacus variant. This isolates two distinct questions that the place-value parameterization combines: (a) does Abacus help if the relevant position embeddings receive gradient (the curriculum case), and (b) does Abacus help *without* a curriculum, just by virtue of its parameterization?


The code, configuration, and the full reproducibility pipeline (`make figures`) are open-source.

2. Setup

Model

Hyperparameter	Value
Layers	6
d_model	384
Attention heads	6
d_head	64
d_ff	1536
Activation	GeLU
Normalization	RMSNorm (pre-norm)
Total parameters	~10.66M

Causal self-attention. The output head is untied from the input embedding. Default `max_len = 32`, set by `max_len_for("addition", 6)` so the model can ingest sequences up to 6-digit-operand addition without architectural changes between in- and out-of-distribution evaluation.



Positional information enters at one of two sites: (a) **added to token embeddings** before the first decoder block (`baseline`, `reversed`, and both Abacus variants), or (b) **applied to Q and K** inside each decoder block's attention computation (`rope`). The `nope` variant skips both. This means RoPE alone is invisible at the embedding stage but appears inside every attention layer; the other PE schemes are localized to the single addition step shown above.

Tokenization

Character-level vocabulary, 15 tokens total:

```
<pad> '0' '1' '2' '3' '4' '5' '6' '7' '8' '9' ' ' '+' '*' '='
```

The longest possible 3-digit-addition example is "999 + 999 = 1998" (16 characters); the longest 6-digit example is "999999 + 999999 = 1999998" (25 characters).

Training

- **Optimizer:** AdamW ($\beta_1 = 0.9$, $\beta_2 = 0.95$, weight decay 0.01).
- **Schedule:** 200-step linear warmup, cosine decay to 10% of peak.
- **Peak LR:** 3×10^{-4} .
- **Batch size:** 512.
- **Epochs:** 8.
- **Loss:** masked cross-entropy on answer tokens *and one terminal pad slot* (the latter teaches the model to stop). Loss on prompt tokens and on pad tokens beyond the first is masked out.

Evaluation

For each digit count $d \in \{1, \dots, 6\}$ I sample 500 random (a, b) pairs with both operands having exactly d digits (operands in $[10^{(d-1)}, 10^d)$, or $[0, 10)$ for $d=1$). Greedy decoding from the prompt "a + b = ", exact-match on the entire decoded answer string. Per-digit-position correctness is also collected over the same sample to feed the heatmap of Section 6.

3. Positional encoding variants

Architecture, optimizer, schedule, and data are held identical across variants. The single axis of variation is how (and whether) position information enters the model.

Variant	Positional encoding	Answer order	Training data
baseline	learned absolute embedding	natural (579)	≤ 3 -digit
reversed	learned absolute embedding	reversed (975)	≤ 3 -digit
nope	none	reversed	≤ 3 -digit
rope	rotary applied to Q, K	reversed	≤ 3 -digit
abacus	digit-position embedding only	reversed	≤ 3 -digit
abacus_curriculum	digit-position embedding only	reversed	mixed 1..5-digit

3.1 Learned absolute position embedding

The standard variant from the original transformer: a learned `Embed(max_len, d_model)` table whose row k is added to the token embedding at sequence position k . Two flavors:

- `baseline`: natural answer order. The model emits the most significant digit first.
- `reversed`: the answer is reversed at training time, so the model emits the ones digit first. Carries flow left-to-right, matching the causal direction.

The reversed flavor is the well-known “Nogueira trick” [Nogueira et al. 2021](#) and is included for two reasons. First, it isolates whether the gain from `reversed` over `baseline` comes from answer order (the carry-friendly direction) or from some other interaction. Second, it provides a fair baseline against which to evaluate the other reversed-answer variants (`nope`, `rope`, `abacus`).

3.2 NoPE

No positional encoding is added at all. The model relies only on the causal-mask asymmetry to break symmetry between positions. Despite this seeming impoverishment, NoPE has been shown [Kazemnejad et al. 2023](#) to extrapolate further than learned absolute PE in several settings, presumably because there is no out-of-distribution position to be confused by.

3.3 RoPE

Rotary position embeddings [Su et al. 2021](#). I precompute a $(T, d_{\text{head}}/2)$ table of $\cos/\sin$ angles $\text{pos} \cdot 10000^{(-2k/d_{\text{head}})}$ and rotate pairs of feature dimensions of Q and K inside `CausalSelfAttention` before computing attention scores. V is not rotated. This formulation expresses position information through *relative* offsets in the attention dot product, which in principle generalizes to unseen absolute positions as long as the relative offsets stay in distribution.

3.4 Abacus place-value embeddings

Following [McLeish et al. 2024](#), absolute position information is dropped entirely and replaced with an embedding of *each digit token’s place value*. The procedure:

1. Identify each contiguous run of digit tokens in the input.
2. For the first two runs (operands, MSB-first by convention), assign position `run_length - index` so the ones digit gets position 1, tens get 2, hundreds get 3, etc.
3. For the third run (the answer, LSB-first under reversed answers), assign position `index + 1` directly — the leftmost digit of the reversed string is the ones digit.
4. Non-digit tokens (including pad, space, +, =) receive position 0, a distinct “no-position” embedding.

The full algorithm in code (from `src/addition_transformer/data.py`):

```
def compute_digit_positions(input_ids, *, reverse_answer):
    """Per-token place-value index. 0 = non-digit; 1+ = distance from ones place."""
```

```

is_digit = (input_ids >= 1) & (input_ids <= 10)  # vocab IDs 1..10 are '0'..'9'
positions = np.zeros_like(input_ids, dtype=np.int32)

# Find each contiguous run of digit tokens.
runs, in_run, start = [], False, 0
for i, d in enumerate(is_digit):
    if d and not in_run:
        start, in_run = i, True
    elif not d and in_run:
        runs.append((start, i))
        in_run = False
if in_run:
    runs.append((start, len(is_digit)))

# Operands (first two runs) are MSB-first; the answer (third run) is LSB-first
# under reversed answers. Both schemes count the same place value (ones=1, tens=2,
# ...), they just read it off different ends of the run.
for run_idx, (s, e) in enumerate(runs):
    run_len = e - s
    is_answer = run_idx == 2
    for j in range(run_len):
        if is_answer and reverse_answer:
            positions[s + j] = j + 1  # LSB-first
        else:
            positions[s + j] = run_len - j  # MSB-first
return positions

```

A `(max_digit_pos + 1, d_model)` learned table is added at every position in the input. With `max_digit_pos = 16` (capping at 16-digit place values), the additional parameter count is negligible relative to the 10.6M total.

Critically, the place-value embedding for position k receives gradient **only if a training example contains a digit at place value k** . For 3-digit-operand training:

- Positions 1, 2, 3 (ones, tens, hundreds) appear in every operand. ✓
- Position 4 (thousands) appears whenever the sum carries, e.g. $999 + 999 = 1998$. ✓
- Positions 5, 6, ..., 16 **never appear**. ✗

This produces the same failure mode as learned absolute PE — embeddings for unseen places stay at initialization — but localized differently in the table. The next variant addresses this.

3.5 Abacus with a length curriculum

Same architecture and embedding scheme as `abacus`, but the training data is sampled uniformly across operand digit counts in $[1, 5]$ rather than fixed at 3 digits. With this data distribution, every place-value embedding from position 1 to position 6 (5-digit operands plus carry) receives gradient during training. The model is then evaluated at digit counts up to 6, requiring it to extrapolate by exactly one place value.

This variant separates two concerns the McLeish et al. result implicitly bundles: (a) the parameterization (place value rather than absolute position) and (b) the training distribution covering the relevant positions. If `abacus_curriculum` generalizes substantially better than `abacus`, then the parameterization advantage is real but conditional on adequate data coverage.

4. The embedding-drift diagnostic

Most discussion of “positions stay at random init” in the length-generalization literature appears as a *plausibility argument*: gradient cannot reach embedding rows associated with positions that the model never has to predict from. I promote this to a directly measurable quantity.

Define, for each row k of a positional embedding table E :

```
drift_k = || E_trained[k, :] - E_init[k, :] ||2
```

where `E_init` is computed by re-instantiating the model with the same RNG seed used to train it. Because JAX initialization is deterministic at fixed seed, `E_init` is bit-identical to the trained model’s initial state. A row that received zero gradient over training has `drift_k = 0.0` exactly.

The diagnostic applies to two tables:

- `pos_emb` (learned absolute PE): drift should be > 0 for positions reachable from the loss-masked region of the longest training example, and exactly 0 for higher positions.
- `digit_pos_emb` (Abacus): drift should be > 0 for place values present in the training data, and exactly 0 beyond.

The diagnostic is cheap (a single L2 norm per row of one or two embeddings) and is stored alongside per-variant accuracy in `results/sweep.json`. The accompanying figure (`results/embedding_drift.png`) is, in my view, the most direct evidence of the failure mode that the length-generalization literature has so far argued for indirectly.

A 20-step CPU smoke run reproduces the expected pattern: `pos_emb` positions 16–25 have drift exactly 0.0, and `digit_pos_emb` positions 5–16 of the clean Abacus variant have drift exactly 0.0, while position 4 — the carry into thousands when summing 3-digit numbers — has drift > 0 as expected.

5. Experiments

For each variant the protocol is: (i) train for 8 epochs on the appropriate dataset (3-digit grid for non-curriculum variants; 1M mixed-digit samples for the curriculum); (ii) evaluate exact-match accuracy at digit counts 1 through 6 with 500 sampled pairs each; (iii) compute per-digit-position accuracy over the same sample; and (iv) record per-row drift for any learned positional embedding present in the model.

The full sweep produces a single `results/sweep.json` that drives three figures:


- **Figure 1 (length_gen.png)** — exact-match accuracy vs operand digit count, one curve per variant. The headline result.
- **Figure 2 (per_digit_heatmap.png)** — per-digit-position error rate, one panel per variant. Reveals whether errors cluster at the most-significant end (typical) or are uniformly distributed (which would indicate a more fundamental failure).
- **Figure 3 (embedding_drift.png)** — per-row L2 drift from init, separated into a learned-absolute panel and an Abacus panel. Reveals which rows actually received gradient.

All artifacts are reproducible from a single command (`make figures`) on a GPU box, or via the included Colab notebook on a free T4.

6. Results

The sweep ran on a single CPU for 12235 seconds (~3.4 hours) with 4 epochs per variant per seed on a 100K-sample subset of the 3-digit grid (50K mixed-digit samples for the curriculum variant). Three random seeds; numbers below are mean $\pm$ standard deviation across seeds. Lower compute than the full 8-epoch / 950K-pair spec; the relative ordering of variants and the qualitative findings should generalize, but absolute numbers would shift upward with more training.

6.1 Exact-match accuracy




Length-generalization curves across all six variants. Shaded bands are $\pm 1\sigma$ over 3 seeds. The gray vertical region marks the training distribution.

Variant	1d	2d	3d	4d	5d	6d
baseline	15.3 $\pm$ 5.8	55.3 $\pm$ 15.6	100.0	0.0	0.0	0.0
reversed	29.7 $\pm$ 26.8	72.0 $\pm$ 36.5	100.0	0.0	0.0	0.0
nope	7.3 $\pm$ 6.9	41.3 $\pm$ 14.6	87.2 $\pm$ 6.9	0.0	0.0	0.0
rope	45.3 $\pm$ 10.6	98.0 $\pm$ 1.4	100.0	0.0	0.0	0.0
abacus	76.7 $\pm$ 17.2	100.0	100.0	0.0	0.0	0.0
abacus_curriculum	100.0	100.0	100.0	100.0	100.0	0.0

Four observations:

- **The 3 $\rightarrow$ 4 digit cliff is the most robust finding in the sweep.** Every non-curriculum variant hits exactly 0.0% at 4, 5, and 6 digits across all three seeds — zero variance, zero exceptions. This is *not* a “near-zero with noise” finding; it is mechanistically 0. The drift figure (§6.3) shows why.
- **abacus_curriculum is the only variant that generalizes.** 100.0 $\pm$ 0.0 out to 5 digits (the training maximum), then 0.0 at 6. The curriculum trains every place-value embedding it can reach; the one position past it fails completely.
- **Seed variability is concentrated at the OOD-distribution edges.** reversed swings from 29.7 $\pm$ 26.8 at 1-digit to 72.0 $\pm$ 36.5 at 2-digit — standard deviations larger than half the mean. By contrast, in-distribution 3-digit accuracy is 100.0 $\pm$ 0.0 for everything except NoPE. A single-seed run would report point estimates for the intermediate cells that land roughly one standard deviation off the mean — which is why the multi-seed protocol matters precisely here, and nowhere else.
- **NoPE is harder to train consistently than the other variants.** Its in-distribution 3-digit accuracy is 87.2 $\pm$ 6.9% — the only variant whose 3-digit numbers are below ceiling, with one of the three seeds plateauing at 77.5%. This is a previously-unstated finding from the multi-seed run: NoPE not only fails to extrapolate, it is *also* an outlier in training stability at this scale and compute budget.

6.2 Per-digit-position pattern




Per-digit-position error rate. Rows are operand digit count; columns are position within the answer (0 = ones place). White = always correct, dark red = always wrong; cells outside the staircase have no support (answer too short to have that position).

The qualitative difference between clean Abacus and every other variant is the headline of this figure. For `baseline`, `reversed`, `nope`, and `rope`, the out-of-distribution region (rows 4–6) is uniformly dark red — *every* digit position is wrong roughly all the time.

The `abacus` (clean) panel is dramatically different. Within the out-of-distribution rows, error rate grades from white at the ones position to dark red at the most-significant position, with the boundary tracking the training-distribution limit (place values 1–4 are correct; place value 5+ fail). Concretely, at 6-digit operands the clean Abacus model gets the ones digit right **98%** of the time, the tens digit **91%**, the hundreds digit **74%**, and then collapses: thousands **6%**, ten-thousands **0%**, hundred-thousands **0%**. The model has correctly learned ones/tens/hundreds addition and applies it at any sequence length; the failure is entirely positional and entirely confined to place values whose embedding rows were never touched by gradient.

`abacus_curriculum` is the cleanest panel: white nearly everywhere, with a small dark band at position 5 and a full-dark column at position 6 of the 6-digit row — the exact two positions just past the curriculum's training maximum.

6.3 Embedding drift



Per-row L2 distance from initialization, averaged over 3 seeds with $\pm 1\sigma$ error bars. Top panel: learned absolute PE. Bottom panel: Abacus digit-position embedding. The brick wall at the training-distribution boundary in each panel is the failure mode the paper is about, and it holds across seeds with very small bars.

Mean drift values across the three seeds:

- **baseline.pos_emb**: positions 0–15 drift in $[0.05, 0.30]$; positions 16–25 drift in $[0.001, 0.0013]$. A $\sim 200\text{--}250\times$ drop at the boundary, and the boundary itself is identical across all three seeds (positions 16+ are below the 0.005 threshold for every seed).
- **reversed.pos_emb**: positions 0–15 drift in $[0.05, 0.21]$; positions 16–25 below 0.005. Same boundary.
- **abacus.digit_pos_emb**: positions 0–4 drift in $[0.03, 0.37]$; positions 5–16 below 0.005. Position 4 — the *carry into thousands*, which appears when summing two 3-digit numbers — has the highest drift (mean 0.37, std ~ 0.05), confirming that even in 3-digit-only training this place value is exercised by the loss.

- **abacus_curriculum.digit_pos_emb**: positions 0–6 drift in [0.03, 0.22]; positions 7–16 below 0.005. Position 6 — the carry from 5-digit additions — has the highest drift, again matching the carry-out-of-training-range expectation.

The drift cliff is bit-exact across seeds. Across all three random initializations and all four embedding tables shown above, the boundary between “received gradient” and “essentially untouched” lands at the same index every time. This is the strongest reproducibility result in the paper.

The ~ 0.001 floor at untouched positions corresponds to the multiplicative shrink from AdamW weight decay over the training run ($\sim 740 \text{ steps} \times 3 \times 10^{-4} \text{ LR} \times 0.01 \text{ weight_decay} \approx 2 \times 10^{-3}$ relative drift, matching the observed value). It is *not* gradient signal; it is a known artifact of weight decay applying to every parameter regardless of gradient. The diagnostic could be sharpened by using cosine drift instead of L2 (cosine is invariant to multiplicative rescaling), but at a 100–300 $\times$ separation the L2 metric is already unambiguous in identifying which positions are functionally untouched.

7. Discussion

The three figures, read together, tell a sharper story than the introduction forecast:

1. **The failure of learned absolute PE at length generalization is mechanistic, not statistical.** It is not that the model “fails to learn” the higher positions; it is that gradient barely reaches those embedding rows ($\sim 100\text{--}300\times$ less drift than trained positions, where the residual is just weight decay). The drift figure makes this directly visible: positions 16–25 of `pos_emb` are a flat ribbon of near-init values, then the bars stop dead at position 15.
2. **Reparameterization alone is not enough — even when the parameterization is “right.”** Clean Abacus, despite its place-value embedding being exactly the structure that *should* help, cliffs from 100% to 0% exact-match at 4 digits, identically to learned PE. The per-position heatmap shows what is happening: ones, tens, and hundreds are nearly always right at OOD digit counts (the model has correctly learned arithmetic at those place values and can apply it at any sequence length), but thousands-and-up are at near- random because their embeddings are at init. Reparameterization gives you *partial* correctness for free; it does not give you generalization unless the training data also covers the relevant positions.
3. **At $\sim 10\text{M}$ parameters and this compute budget, neither NoPE nor RoPE extrapolates.** Both cliff at the 3 $\rightarrow$ 4-digit boundary just as hard as learned absolute PE. This contradicts a common framing in the length-generalization literature (“NoPE > learned PE for extrapolation”, “RoPE has unlimited extrapolation”). Two caveats: (a) those results are typically reported at much larger scale and on natural language, where extrapolation may rely on different mechanisms; (b) the compute budget here is modest (4 epochs on 100K samples per variant). It is possible RoPE and NoPE would extrapolate further with more training, though the drift evidence suggests that for *arithmetic specifically*, position-agnostic or relative-position schemes still need to see each place value.

4. **The curriculum + Abacus combination is the only variant that generalizes.**

`abacus_curriculum` is 100% accurate at 4 and 5 digits despite a training distribution that includes those sizes — and then cliffs at 6, exactly one position past the curriculum’s max. The drift diagnostic confirms this is not a mystery: position 6’s place-value embedding is the largest-drift row in the entire model (carry from 5-digit additions), while position 7 is at near-init. The curriculum variant extrapolates by exactly one position, and only one.

5. **Per-digit errors are mostly carries.** The lowest-order digit of a sum is computable from the low-order digits of the operands alone, with no carry propagation. The highest-order digit requires the entire carry chain to be correct. The heatmap therefore grades red from right (low-order) to left (high-order), and the boundary creeping leftward as digit count grows is essentially a picture of how far the model can propagate carries. The clean-Abacus heatmap makes this especially vivid: it has the cleanest left-to-right gradient because nothing else is in the way — no positional confusion, just a hard wall at the highest trained place value.

6. **The multi-seed run reveals where the original draft’s point estimates were misleading.**

The 3→4 cliff and the drift boundary are reproducible across seeds to the bit. But the intermediate-digit numbers — especially `reversed` at 1- and 2-digit operands — vary by tens of percentage points across seeds. The standard deviations in §6.1 should be read as *warnings* on any narrative that distinguishes variants by single-digit percentages: at this scale and compute budget, `reversed` could equally well be reported as “9% accurate at 1-digit” or “56% accurate at 1-digit” depending on the seed. The features that survive seed variation — the 3-digit ceiling, the 4-digit cliff, the drift boundary, and the Abacus-curriculum extension to 5 digits — are the load-bearing findings of the paper.

7. **NoPE has an additional, separate failure: convergence stability.** Its in-distribution 3-digit accuracy is $87.2 \pm 6.9\%$, with one of three seeds plateauing at 77.5%. Every other variant hits 100.0 ± 0.0 in distribution. This is *not* a length-generalization story; it is a separate observation that the inductive bias from causal masking alone is enough to *sometimes* converge to a working arithmetic circuit but not always. Practical NoPE deployments at this scale would need either more compute, multiple seed reruns, or a small amount of positional information to be reliable.

8. Limitations

This is a small-scale, single-task, single-seed study with a constrained compute budget. Specifically:

- **Scale.** 10.6M parameters is a few orders of magnitude smaller than typical research models. Results should not be assumed to transfer to billion-parameter models without verification, particularly for NoPE and RoPE, where the inductive-bias arguments may behave differently with greater capacity.
- **Compute.** The sweep ran on a CPU for ~3.4 hours total (3 seeds × 6 variants × 4 epochs on a 100K-sample subset of the 3-digit grid; 50K for the mixed-digit curriculum). The full spec is 8 epochs on the ~950K grid; absolute accuracies would shift upward with more training, and it remains possible that NoPE and RoPE would extrapolate non-trivially given enough optimization time. The qualitative ordering of variants and the drift-cliff finding are bit-exact across seeds and expected to be robust.

- **Task.** Addition is much easier than multiplication, and multiplication is much easier than (say) modular exponentiation. Multiplication is supported in the codebase but addition is the headline task because the failure modes are cleaner.
- **Single-digit eval is OOD.** A setup artifact: training data is uniform over $[0, 1000)$, so single-digit operand pairs make up $\sim 1\%$ of the distribution and the non-curriculum variants effectively never see $5 + 7 =$. The 1-digit column of Table 6.1 measures generalization to a *shorter* but *unseen* length, not in-distribution accuracy. Future iterations should either stratify training data by digit count or report evaluations only at digit counts that the variant explicitly saw.
- **Seeds.** The sweep uses 3 random seeds, which is enough to give meaningful error bars on the cliff-boundary comparison but not enough to detect rarer failure modes (the NoPE convergence instability shows up as 1-of-3 here, so the actual rate could be anywhere from $\sim 10\%$ to $\sim 70\%$ with reasonable confidence). 5–10 seeds would tighten the variance estimates further.
- **Decoding.** All evaluation uses greedy decoding. Beam search or temperature sampling may shift accuracy upward, particularly at the boundary of generalization.
- **Curriculum design.** The length curriculum used here is the simplest possible (uniform mixture of digit counts 1–5). Curriculum design itself is a research question I do not engage with here.

9. Related work

Positional encodings as the proximate cause of failure. [Vaswani et al. 2017](#) introduced sinusoidal absolute position embeddings, later commonly replaced by learned absolute PE for simplicity. [Su et al. 2021](#) introduced rotary position embeddings (RoPE), expressing position through rotations of Q and K and claiming better extrapolation properties. [Press et al. 2022](#) introduced ALiBi, a static linear bias on attention scores — not evaluated here, but a natural fifth point of comparison. [Kazemnejad et al. 2023](#) argued that for decoder-only models, **no positional encoding at all** can extrapolate further than learned absolute PE, attributing this to the implicit position information carried by causal masking. The results in §6 do not reproduce this advantage at 10M-parameter scale and the compute budget here; whether the gap closes at larger scale or with more training is left to future work.

Arithmetic-specific interventions. [Nogueira et al. 2021](#) observed that reversing the answer (predict ones digit first) makes carries flow with the causal direction, the trick used in every reversed-answer variant here. [Nye et al. 2021](#) introduced scratchpad reasoning, where the model emits intermediate carry computations; this trades inference compute for accuracy. [Lee et al. 2023](#) systematically studied small transformers on arithmetic and identified the reversed-answer trick as critical for in-distribution accuracy at 3+ digits. [McLeish et al. 2024](#) introduced Abacus embeddings — a learned table indexed by a digit’s distance from the ones place — and reported strong length-generalization gains. The clean-Abacus / Abacus-curriculum comparison in this report is designed to disentangle the parameterization from the data-coverage component of that result; §6 shows that the parameterization alone does not suffice.

Length generalization broadly. [Anil et al. 2022](#) demonstrated that length generalization is qualitatively hard for vanilla transformers even with perfect in-distribution accuracy. The drift-cliff finding in §6.3 gives one mechanistic answer: at the transition out of distribution, embedding rows have not received any gradient signal and therefore cannot encode position-conditional information. [Power et al. 2022](#) showed

that small transformers on modular arithmetic exhibit *grokking* — abrupt generalization long after training loss saturates. This study uses non-modular addition and does not focus on grokking, but the underlying optimization dynamics are related.

10. Conclusion

This report presents a controlled comparison of six positional-encoding schemes on a 10M-parameter arithmetic transformer, and introduces an embedding-drift diagnostic that makes the dominant length-generalization failure mode directly measurable in the trained weights. The headline finding is that at this scale, *every* mechanism-only intervention — NoPE, RoPE, clean Abacus — cliffs at the training-distribution boundary just as hard as learned absolute PE. The only variant that generalizes past 3 digits is Abacus combined with a length curriculum that exposes every relevant place value to gradient, and even this variant extrapolates by exactly one position before cliffing in turn. The drift diagnostic confirms in the weights what the accuracy curves show in the outputs: rows of the positional embedding that never received gradient stay within a weight-decay-only band of initialization, and accuracy collapses precisely at the positions where the drift collapses.

The methodological contribution — turning a plausibility argument about gradients into a number computed from the trained weights — appears to me more durable than the specific empirical finding, which is constrained by scale, single-seed evaluation, and a constrained compute budget. The diagnostic is two lines of code and should apply unchanged to any model with an addressable positional-embedding table.

Future work includes (i) replicating across multiple seeds with full training, (ii) extending to multiplication and modular operations, (iii) overlaying attention-pattern visualizations on the partial-correctness pattern of clean Abacus, and (iv) probing whether NoPE and RoPE actually extrapolate at larger scales, since the failure here may be a compute-budget artifact rather than a fundamental ceiling.

References

1. Anil, C. et al. (2022). *Exploring length generalization in large language models*. NeurIPS. [arXiv:2207.04901](#)
2. Kazemnejad, A. et al. (2023). *The Impact of Positional Encoding on Length Generalization in Transformers*. NeurIPS. [arXiv:2305.19466](#)
3. Lee, N. et al. (2023). *Teaching Arithmetic to Small Transformers*. [arXiv:2307.03381](#)
4. McLeish, S. et al. (2024). *Transformers Can Do Arithmetic with the Right Embeddings*. NeurIPS. [arXiv:2405.17399](#)
5. Nogueira, R., Jiang, Z., Lin, J. (2021). *Investigating the Limitations of Transformers with Simple Arithmetic Tasks*. [arXiv:2102.13019](#)
6. Nye, M. et al. (2021). *Show Your Work: Scratchpads for Intermediate Computation with Language Models*. [arXiv:2112.00114](#)

7. Power, A. et al. (2022). *Grokking: Generalization Beyond Overfitting on Small Algorithmic Datasets*. [arXiv:2201.02177](#)
8. Press, O., Smith, N., Lewis, M. (2022). *Train Short, Test Long: Attention with Linear Biases Enables Input Length Extrapolation*. ICLR. [arXiv:2108.12409](#)
9. Su, J. et al. (2021). *RoFormer: Enhanced Transformer with Rotary Position Embedding*. [arXiv:2104.09864](#)
10. Vaswani, A. et al. (2017). *Attention Is All You Need*. NeurIPS. [arXiv:1706.03762](#)